

Md.Tanbin Emon

In [11]:

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

In [12]:

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import classification_report, accuracy_score
```

In [13]:

```
data = pd.read_csv("/content/drive/MyDrive/eadg /Final/Sample_Data_AI_Lab_Final.xlsx - Sheet1.csv")
```

Handle Missing Values

In [15]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5986 entries, 0 to 5985
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   MultipleLines          5986 non-null   int64
1   InternetService        5986 non-null   int64
2   OnlineSecurity         5986 non-null   int64
3   OnlineBackup           5986 non-null   int64
4   DeviceProtection       5986 non-null   int64
5   TechSupport            5986 non-null   int64
6   StreamingTV            5986 non-null   int64
7   StreamingMovies        5986 non-null   int64
8   MonthlyCharges         5986 non-null   float64
9   TotalCharges           5986 non-null   object
10  Contract               5986 non-null   object
dtypes: float64(1), int64(8), object(2)
memory usage: 514.5+ KB
```

In [16]:

```
data.columns
```

Out[16]:

```
Index(['MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup',
      'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies',
      'MonthlyCharges', 'TotalCharges', 'Contract'],
      dtype='object')
```

In [17]:

```
data.isnull().sum()
```

Out[17]:

	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
MonthlyCharges	0
TotalCharges	0
Contract	0

dtype: int64

In [18]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

Encoding Categorical Variables

In [31]:

```
# Apply One-Hot Encoding
data_one_hot = pd.get_dummies(data, columns=['Contract'], drop_first=True)

# Check the transformed data
print(data_one_hot.head())
```

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	\
0	1	0	3	3	
1	0	2	0	1	
2	1	2	0	0	
3	0	1	0	0	
4	0	1	1	0	

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	\
0	3	3	3	3	
1	1	0	1	0	
2	0	0	0	0	
3	0	0	0	1	
4	1	0	0	0	

	MonthlyCharges	TotalCharges	Contract_0.5	Contract_1.0
0	24.10	1734.65	False	True
1	88.15	3973.2	False	False
2	74.95	2869.85	False	False
3	55.90	238.5	False	False
4	53.45	119.5	False	False

In [51]:

```
data
```

Out[51]:

0	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMo
1	0	2	0	1	1	0	1	
2	1	2	0	0	0	0	0	
3	0	1	0	0	0	0	0	
4	0	1	1	0	1	0	0	
...	
5981	0	2	1	0	0	0	1	
5982	1	1	1	1	1	1	1	
5983	0	0	3	3	3	3	3	
5984	1	2	0	0	1	0	1	
5985	0	0	3	3	3	3	3	

5986 rows x 11 columns




In [20]:

```
data
```

Out[20]:

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMo
0	1	0	3	3	3	3	3	
1	0	2	0	1	1	0	1	
2	1	2	0	0	0	0	0	
3	0	1	0	0	0	0	0	
4	0	1	1	0	1	0	0	
...	
5981	0	2	1	0	0	0	1	
5982	1	1	1	1	1	1	1	
5983	0	0	3	3	3	3	3	
5984	1	2	0	0	1	0	1	
5985	0	0	3	3	3	3	3	

5986 rows x 11 columns



Feature Scaling

In [32]:

```

from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()

data['Contract'] = label_encoder.fit_transform(data['Contract'])

from sklearn.preprocessing import MinMaxScaler

min_max_scaler = MinMaxScaler()
contract_values = data['Contract'].values.reshape(-1, 1)
data['Contract'] = min_max_scaler.fit_transform(contract_values)

print(data['Contract'].head())

```

```

0      1.0
1      0.0
2      0.0
3      0.0
4      0.0
Name: Contract, dtype: float64

```

In [33]:

```
data
```

Out[33]:

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMo
0	1	0	3	3	3	3	3	
1	0	2	0	1	1	0	1	
2	1	2	0	0	0	0	0	
3	0	1	0	0	0	0	0	
4	0	1	1	0	1	0	0	
...
5981	0	2	1	0	0	0	1	
5982	1	1	1	1	1	1	1	
5983	0	0	3	3	3	3	3	
5984	1	2	0	0	1	0	1	
5985	0	0	3	3	3	3	3	

5986 rows x 11 columns



Dataset Splitting

In [34]:

```

from sklearn.model_selection import train_test_split

X = data[['MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup',
          'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies',
          'MonthlyCharges', 'TotalCharges']]
y = data['Contract']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Display the shapes of the resulting datasets
print("Training Features Shape:", X_train.shape)
print("Testing Features Shape:", X_test.shape)
print("Training Target Shape:", y_train.shape)

```

```
print("Testing Target Shape:", y_test.shape)
```

Training Features Shape: (4788, 10)
Testing Features Shape: (1198, 10)
Training Target Shape: (4788,)
Testing Target Shape: (1198,)


In [27]:

data

Out[27]:

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMo
0	1	0	3	3	3	3	3	
1	0	2	0	1	1	0	1	
2	1	2	0	0	0	0	0	
3	0	1	0	0	0	0	0	
4	0	1	1	0	1	0	0	
...
5981	0	2	1	0	0	0	1	
5982	1	1	1	1	1	1	1	
5983	0	0	3	3	3	3	3	
5984	1	2	0	0	1	0	1	
5985	0	0	3	3	3	3	3	

5986 rows x 11 columns



In [35]:

```
print("Training Features Shape:", X_train)
print("Testing Features Shape:", X_test)
print("Training Target Shape:", y_train)
print("Testing Target Shape:", y_test)
```

Training Features Shape:	MultipleLines	InternetService	OnlineSecurity	OnlineBack	up \
4779	1	2	0	1	
685	1	2	0	1	
2993	1	1	1	0	
3146	1	1	1	1	
35	3	1	0	1	
...	
5876	1	2	0	0	
4093	3	1	0	1	
2496	1	2	0	1	
4163	1	2	1	0	
2572	0	2	1	0	
DeviceProtection	TechSupport	StreamingTV	StreamingMovies	\	
4779	0	0	1	1	
685	1	0	1	1	
2993	1	0	1	1	
3146	1	1	0	1	
35	0	1	0	0	
...	
5876	1	1	0	0	
4093	1	0	1	1	
2496	0	0	0	0	
4163	0	0	1	1	
2572	0	0	0	0	
MonthlyCharges	TotalCharges				
4779	0.0 0.0	4744	25		

```
4779      99.00      4744.33
685      104.75      5510.65
2993      79.70      5293.4
3146      79.90      3326.2
35      35.65      425.1
...
5876      86.75      1410.25
4093      55.45      2966.95
2496      79.15      79.15
4163      99.60      347.65
2572      73.60      73.6
```

[4788 rows x 10 columns]

```
Testing Features Shape:      MultipleLines  InternetService  OnlineSecurity  OnlineBacku
p \
3997      0      1      1      0
1147      0      1      1      0
3659      0      0      3      3
2951      0      2      0      0
150      1      1      0      1
...
4880      1      2      0      1
1903      0      2      0      0
5507      0      1      1      1
5069      3      1      0      1
3574      1      2      0      1
```

```
      DeviceProtection  TechSupport  StreamingTV  StreamingMovies  \
3997      1      0      1      0
1147      1      0      1      0
3659      3      3      3      3
2951      0      0      1      0
150      1      0      1      0
...
4880      0      0      0      0
1903      0      1      1      1
5507      1      1      1      1
5069      1      0      1      0
3574      1      0      1      1
```

```
      MonthlyCharges  TotalCharges
3997      64.25      1024
1147      66.40      94.55
3659      20.15      1337.5
2951      78.90      299.75
150      68.95      2038.7
...
4880      80.65      1451.9
1903      94.45      1511.2
5507      85.00      5607.75
5069      46.40      812.4
3574      104.80      7470.1
```

[1198 rows x 10 columns]

Training Target Shape: 4779 0.5

```
685      0.5
2993      1.0
3146      0.0
35      0.0
```

```
...
5876      0.5
4093      0.5
2496      0.0
4163      0.0
2572      0.0
```

Name: Contract, Length: 4788, dtype: float64

Testing Target Shape: 3997 1.0

```
1147      0.0
3659      1.0
2951      0.0
150      0.5
```

```
...
4000      0.0
```

```
4880      0.0
1903      0.0
5507      1.0
5069      0.5
3574      1.0
Name: Contract, Length: 1198, dtype: float64
```

Classification Task Implement the following classification algorithms to predict the target variable

- K-Nearest Neighbors (KNN)
- Random Forest Classifier
- Naïve Bayes Classifier

In [28]:


```
from sklearn.datasets import make_blobs
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
```


In [40]:

```
X, y = make_blobs(n_samples = 500, n_features = 2, centers = 4, cluster_std = 1.5, random_state = 41)
```

In [41]:

```
sns.set_theme()
plt.figure(figsize=(10, 10))
plt.scatter(X[:, 0], X[:, 1], c=y, marker='*', s=100, edgecolors='black')
plt.show()
```





In [42]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 0)
```

In [43]:

```
knn5 = KNeighborsClassifier(n_neighbors = 5)
knn1 = KNeighborsClassifier(n_neighbors=1)
```

In [44]:

```
knn5.fit(X_train, y_train)
knn1.fit(X_train, y_train)

y_pred_5 = knn5.predict(X_test)
y_pred_1 = knn1.predict(X_test)
```

In [45]:

```
from sklearn.metrics import accuracy_score
print("Accuracy with k=5", accuracy_score(y_test, y_pred_5)*100)
print("Accuracy with k=1", accuracy_score(y_test, y_pred_1)*100)
```

Accuracy with k=5 83.2
Accuracy with k=1 83.2

KNN

In [46]:

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report

# Initialize and train the KNN model
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

# Predict using the KNN model
y_pred_knn = knn.predict(X_test)

# Generate and print the classification report
print("K-Nearest Neighbors (KNN) Classification Report:")
print(classification_report(y_test, y_pred_knn))
```

K-Nearest Neighbors (KNN) Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	33
1	1.00	1.00	1.00	34
2	0.66	0.68	0.67	31
3	0.62	0.59	0.60	27
accuracy			0.83	125
macro avg	0.82	0.82	0.82	125
weighted avg	0.83	0.83	0.83	125

Random Forest Classifier

In [47]:

```
from sklearn.ensemble import RandomForestClassifier
```



```
# Initialize and train the Random Forest model
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# Predict using the Random Forest model
y_pred_rf = rf.predict(X_test)

# Generate and print the classification report
print("Random Forest Classifier Classification Report:")
print(classification_report(y_test, y_pred_rf))
```

```
Random Forest Classifier Classification Report:
              precision    recall  f1-score   support

    0           1.00        1.00        1.00         33
    1           1.00        1.00        1.00         34
    2           0.69        0.65        0.67         31
    3           0.62        0.67        0.64         27

 accuracy          0.84
macro avg          0.83        0.83        0.83        125
weighted avg       0.84        0.84        0.84        125
```

Naïve Bayes Classifier

In [48]:

```
from sklearn.naive_bayes import GaussianNB

# Initialize and train the Naive Bayes model
nb = GaussianNB()
nb.fit(X_train, y_train)

# Predict using the Naive Bayes model
y_pred_nb = nb.predict(X_test)

# Generate and print the classification report
print("Naïve Bayes Classifier Classification Report:")
print(classification_report(y_test, y_pred_nb))
```

```
Naïve Bayes Classifier Classification Report:
              precision    recall  f1-score   support

    0           1.00        1.00        1.00         33
    1           1.00        1.00        1.00         34
    2           0.74        0.74        0.74         31
    3           0.70        0.70        0.70         27

 accuracy          0.87
macro avg          0.86        0.86        0.86        125
weighted avg       0.87        0.87        0.87        125
```

Evaluation Metrics For each classifier, generate a detailed **classification report** that includes the following metrics:

1. Precision
2. Recall
3. F1-Score
4. Accuracy

In [49]:

```
from sklearn.metrics import classification_report, accuracy_score

# Function to print evaluation metrics
def evaluate_model(y_true, y_pred, model_name):
    print(f"--- {model_name} Classification Report ---")
```

```

print(classification_report(y_true, y_pred))
accuracy = accuracy_score(y_true, y_pred)
print(f"Accuracy: {accuracy:.4f}\n")

# K-Nearest Neighbors (KNN)
evaluate_model(y_test, y_pred_knn, "K-Nearest Neighbors (KNN)")

# Random Forest Classifier
evaluate_model(y_test, y_pred_rf, "Random Forest Classifier")

# Naive Bayes Classifier
evaluate_model(y_test, y_pred_nb, "Naive Bayes Classifier")

```

```

--- K-Nearest Neighbors (KNN) Classification Report ---
              precision    recall  f1-score   support

    0           1.00        1.00        1.00         33
    1           1.00        1.00        1.00         34
    2           0.66        0.68        0.67         31
    3           0.62        0.59        0.60         27

 accuracy                   0.83         125
macro avg           0.82        0.82        0.82         125
weighted avg        0.83        0.83        0.83         125

```

Accuracy: 0.8320

```

--- Random Forest Classifier Classification Report ---
              precision    recall  f1-score   support

    0           1.00        1.00        1.00         33
    1           1.00        1.00        1.00         34
    2           0.69        0.65        0.67         31
    3           0.62        0.67        0.64         27

 accuracy                   0.84         125
macro avg           0.83        0.83        0.83         125
weighted avg        0.84        0.84        0.84         125

```

Accuracy: 0.8400

```

--- Naive Bayes Classifier Classification Report ---
              precision    recall  f1-score   support

    0           1.00        1.00        1.00         33
    1           1.00        1.00        1.00         34
    2           0.74        0.74        0.74         31
    3           0.70        0.70        0.70         27

 accuracy                   0.87         125
macro avg           0.86        0.86        0.86         125
weighted avg        0.87        0.87        0.87         125

```

Accuracy: 0.8720

Comparative Analysis Compare the performance of the three algorithms using the metrics generated in the classification reports. Include a brief discussion addressing:

- Which algorithm performs the best based on the metrics.
- The possible reasons for the observed performance differences.
- Recommendations for improving classification performance if necessary.

In [50]:

```

import pandas as pd

metrics = {
    "Model": ["KNN", "Random Forest", "Naive Bayes"],
    "Accuracy": [

```

```

        accuracy_score(y_test, y_pred_knn),
        accuracy_score(y_test, y_pred_rf),
        accuracy_score(y_test, y_pred_nb)
    ],
    "Precision": [
        classification_report(y_test, y_pred_knn, output_dict=True)['weighted avg']['precision'],
        classification_report(y_test, y_pred_rf, output_dict=True)['weighted avg']['precision'],
        classification_report(y_test, y_pred_nb, output_dict=True)['weighted avg']['precision']
    ],
    "Recall": [
        classification_report(y_test, y_pred_knn, output_dict=True)['weighted avg']['recall'],
        classification_report(y_test, y_pred_rf, output_dict=True)['weighted avg']['recall'],
        classification_report(y_test, y_pred_nb, output_dict=True)['weighted avg']['recall']
    ],
    "F1-Score": [
        classification_report(y_test, y_pred_knn, output_dict=True)['weighted avg']['f1-score'],
        classification_report(y_test, y_pred_rf, output_dict=True)['weighted avg']['f1-score'],
        classification_report(y_test, y_pred_nb, output_dict=True)['weighted avg']['f1-score']
    ]
}

metrics_df = pd.DataFrame(metrics)

print("Comparative Analysis of Classification Models:")
print(metrics_df)

print("\nBest Performing Model per Metric:")
print(metrics_df.loc[metrics_df[['Accuracy', 'Precision', 'Recall', 'F1-Score']].idxmax()])

```

Comparative Analysis of Classification Models:

	Model	Accuracy	Precision	Recall	F1-Score
0	KNN	0.832	0.831673	0.832	0.831748
1	Random Forest	0.840	0.841103	0.840	0.840190
2	Naive Bayes	0.872	0.872000	0.872	0.872000

Best Performing Model per Metric:

	Model	Accuracy	Precision	Recall	F1-Score
2	Naive Bayes	0.872	0.872	0.872	0.872
2	Naive Bayes	0.872	0.872	0.872	0.872
2	Naive Bayes	0.872	0.872	0.872	0.872
2	Naive Bayes	0.872	0.872	0.872	0.872